

Libraries

```
In [ ]: # code courtesy of https://www.kaggle.com/alvations/n-gram-language-model-with-nltk

!pip install -U pip
!pip install -U dill
!pip install -U nltk==3.4
```

```
In [ ]: import nltk
from nltk.util import pad_sequence
from nltk.util import bigrams
from nltk.util import ngrams
from nltk.util import everygrams
from nltk.lm.preprocessing import pad_both_ends
from nltk.lm.preprocessing import flatten
from nltk.corpus import reuters
```

For Reutres

```
In [ ]: text =[]
for sentence in reuters.sents():
    text.append(sentence)
```

```
In [ ]: from nltk.util import pad_sequence
list(pad_sequence(text[0],
                 pad_left=True, left_pad_symbol="<s>",
                 pad_right=True, right_pad_symbol="</s>",
                 n=2)) # The n order of n-grams, if it's 2-grams, you pad once, 3-grams pad twice, etc.
```

```
In [ ]: padded_sent = list(pad_sequence(text[0], pad_left=True, left_pad_symbol="<s>",
                                       pad_right=True, right_pad_symbol="</s>", n=2))
list(ngrams(padded_sent, n=2))
```

```
In [ ]: list(pad_sequence(text[0],
                         pad_left=True, left_pad_symbol="<s>",
                         pad_right=True, right_pad_symbol="</s>",
                         n=3)) # The n order of n-grams, if it's 2-grams, you pad once, 3-grams pad twice, etc.
```

```
In [ ]: padded_sent = list(pad_sequence(text[0], pad_left=True, left_pad_symbol="<s>",
                                       pad_right=True, right_pad_symbol="</s>", n=3))
list(ngrams(padded_sent, n=3))
```

Building Language Model

```
In [ ]: from nltk.lm.preprocessing import pad_both_ends
        list(pad_both_ends(text[0], n=2))
```

```
In [ ]: list(bigrams(pad_both_ends(text[0], n=2)))
```

```
In [ ]: from nltk.util import everygrams
        padded_bigrams = list(pad_both_ends(text[0], n=2))
        list(everygrams(padded_bigrams, max_len=2))
```

```
In [ ]: from nltk.lm.preprocessing import flatten
        list(flatten(pad_both_ends(sent, n=2) for sent in text))
```

```
In [ ]: from nltk.lm.preprocessing import padded_everygram_pipeline
        train, vocab = padded_everygram_pipeline(2, text)
```

```
In [ ]: training_ngrams, padded_sentences = padded_everygram_pipeline(2, text)
        for ngramlize_sent in training_ngrams:
            print(list(ngramlize_sent))
            print()
        print('#####')
        list(padded_sentences)
```

For Corpus

```
In [ ]: import os
        import requests
        import io #codecs

        # Text version of https://kilgarriff.co.uk/Publications/2005-K-Lineer.pdf
        if os.path.isfile('language-never-random.txt'):
            with io.open('language-never-random.txt', encoding='utf8') as fin:
                text = fin.read()
        else:
            url = "https://github.com/praveenjoshi01/COMP9066_27375-Natural-Language-P
            rocessing/blob/main/language-never-random.txt"
            text = requests.get(url).content.decode('utf8')
            with io.open('language-never-random.txt', 'w', encoding='utf8') as fout:
                fout.write(text)
```

```
In [ ]: try: # Use the default NLTK tokenizer.
        from nltk import word_tokenize, sent_tokenize
        # Testing whether it works.
        # Sometimes it doesn't work on some machines because of setup issues.
        word_tokenize(sent_tokenize("This is a foobar sentence. Yes it is.")[0])
    except: # Use a naive sentence tokenizer and toktok.
        import re
        from nltk.tokenize import ToktokTokenizer
        # See https://stackoverflow.com/a/25736515/610569
        sent_tokenize = lambda x: re.split(r'(?<=[^A-Z].[.?!]) +(?=[A-Z])', x)
        # Use the toktok tokenizer that requires no dependencies.
        toktok = ToktokTokenizer()
        word_tokenize = word_tokenize = toktok.tokenize
```

```
In [ ]: # Tokenize the text.
        tokenized_text = [list(map(str.lower, word_tokenize(sent)))
                           for sent in sent_tokenize(text)]
```

```
In [ ]: tokenized_text[0]
```

```
In [ ]: # Preprocess the tokenized text for 3-grams language modelling
        n = 3
        train_data, padded_sents = padded_everygram_pipeline(n, tokenized_text)
```

```
In [ ]: from nltk.lm import MLE
        model = MLE(n) # Lets train a 3-grams model, previously we set n=3
```

```
In [ ]: len(model.vocab)
```

```
In [ ]: model.fit(train_data, padded_sents)
        print(model.vocab)
```

```
In [ ]: len(model.vocab)
```

```
In [ ]: model.counts['language']
```

Generation using N-gram Language Model

```
In [ ]: from nltk.tokenize.treebank import TreebankWordDetokenizer

detokenize = TreebankWordDetokenizer().detokenize

def generate_sent(model, num_words, random_seed=42):
    """
    :param model: An ngram language model from `nltk.Lm.model`.
    :param num_words: Max no. of words to generate.
    :param random_seed: Seed value for random.
    """
    content = []
    for token in model.generate(num_words, random_seed=random_seed):
        if token == '<s>':
            continue
        if token == '</s>':
            break
        content.append(token)
    return detokenize(content)
```

```
In [ ]: print(model.generate(28, random_seed=0))
```

```
In [ ]: generate_sent(model, 20, random_seed=42)
```

Computing Perplexity

```
In [ ]: test_sentences = ['he mathematics of the random is well']
tokenized_text = [list(map(str.lower, nltk.tokenize.word_tokenize(sent))) for
sent in test_sentences]

test_data = [nltk.bigrams(t, pad_right=True, pad_left=True, left_pad_symbol="
<s>", right_pad_symbol="</s>") for t in tokenized_text]
for test in test_data:
    print("MLE Estimates:", [(ngram[-1], ngram[:-1]), model.score(ngram[-1],
ngram[:-1])] for ngram in test)

test_data = [nltk.bigrams(t, pad_right=True, pad_left=True, left_pad_symbol="
<s>", right_pad_symbol="</s>") for t in tokenized_text]
for i, test in enumerate(test_data):
    print("PP({0}):{1}".format(test_sentences[i], model.perplexity(test)))
```

```
In [ ]:
```